

TDT4225 Very Large, Distributed Data Volumes

Exercise 4

Deadline 15 Nov 23:59 (midnight)

1. Kleppmann Chap 5

Explain how version numbers are used in the Shopping card example (buying food) when capturing the happens-before relationship to know when to merge in values from “concurrent writes”?

2. Kleppmann Chap 6

- a) What is the best way of supporting re-partitioning? And why is this the best way? (According to Kleppmann).
- b) Explain when you should use local indexing, and when you should use global indexing?

3. Kleppmann Chap 7

- a) **Read committed vs snapshot isolation.** We want to compare *read committed* with *snapshot isolation*. We assume the traditional way of implementing read committed, where write locks are held to the end of the transaction, while read locks are set and released when doing the read itself. Show how the following schedule is executed using these two approaches:

$r_1(A); w_2(A); w_2(B); r_1(B); c_1; c_2;$

- b) Also show how this is executed using serializable with 2PL (two-phase locking).

4. Kleppmann Chap 8


- a) If you send a message in a network and you do not get a reply, what could have happened? List some alternatives.
- b) Explain why and how using clocks for *last write wins* could be dangerous.
- c) Explain how a node sometimes can not trust its own judgement

5. Kleppmann Chap 9


- a) Explain the connection between ordering, linearizability and consensus.
- b) Are there any distributed data systems which are usable even if they are not linearizable? Explain your answer.
- c) What is the main reason for not having a linearizable distributed system?

6. Coulouris Chap 14

- Given two events e and f . Assume that the logical (Lamport) clock values L are such that $L(e) < L(f)$. Can we then deduce that e "happened before" f ? Why? What happens if one uses vector clocks instead? Explain.
- The figure below shows three processes and several events. Vector clock values are given for some of the events. Give the vector clock values for the remaining events.



- The figure below shows the events that occur at two processes P1 and P2. The arrows mean sending of messages. Show the alternative consistent states the system can have had. Start from state S_{00} . (S_{xy} where x is p1's state and y is p2's state)



7. RAFT

RAFT has a concept where the log is replicated to all participants. How does RAFT ensure that the log is equal on all nodes in case of a crash and a new leader?

8. MySQL RAFT

What are the main takeaways (advantages) from introducing RAFT into MySQL?

9. Stonebraker/Pavlo

What is happening with Document vs SQL databases according to Stonebraker/Pavlo?